

In both versions, the environment worked as expected. The agent received percepts such as breeze, stench, glitter, bump, and scream, then selected actions and received rewards from the environment. The random agent often repeated poor actions and eventually entered a dangerous square, which led to a large negative reward.

1. Using only the percepts for training can be a problem because they do not show the full state of the world. They only give local clues about nearby hazards or rewards. The same percept can appear in different situations, so the agent may not know which action is best. This makes the problem partially observable and harder to learn.
2. A table could be used for the value function if we include the full state, because WumpusWorld is small and has a limited number of positions, orientations, and actions. However, if the table only uses percepts, it may not work well because different states can produce the same percept. So a table is possible, but only if the state representation includes enough information.
3. Curriculum training means starting with an easier version of the environment and making it harder over time. For example, there could be fewer pits, no Wumpus, or the gold could be placed closer to the start. This would help the agent learn to explore and collect the gold instead of always leaving immediately. After the agent learns useful behavior, the normal harder environment can be introduced.
4. A problem is Markovian when the current state contains enough information to determine what happens next, so the full past history is not needed. In WumpusWorld, the full environment state is Markovian because it includes the agent's position and direction, the locations of the pits, gold, and Wumpus, and whether the agent still has the arrow or the gold. From that full state, the next state and reward can be determined from the chosen action. However, the agent usually only observes percepts such as breeze, stench, glitter, bump, and scream. These percepts do not reveal the exact state of the world, and the agent may need previous observations to decide what to do. So the full world is Markovian, but from the agent's point of view it is partially observable, which makes it closer to a POMDP.